

Introduction to Julia for Python developers

PyData, Berlin
May, 2016

David Higgins



Who is this guy?



Have been programming and involved in Open Source since the 1980's

Currently working as a computational neuroscientist
-requires both rapid prototyping and scalability

I tend to work on projects which are highly numerical,
requiring both precision and processing power

I've been using Julia in my daily work for 2 years now

And I teach a course on Advanced Scientific
Programming using Python at the Humboldt and
Technical universities

What is Julia?



It is a new programming language (begun 2009, first release 2012)

Produced by a group at MIT with a long history of developing tools for parallelising existing programming languages

It is a General Purpose language with Numerical Computing at the core of its design

Who should be interested in Julia?



Data scientists, anyone interested in performance in numerical computing

If you are already pushing the bounds of Python via Cython

If you want a faster route from prototype to high-performance implementation

If you want a new toy

Why *you* should be interested in Julia



Syntactically it is easy to learn from Python

All Python libraries are fully usable via PyCall
command

The development cycle is similar to Python (rapid
prototyping)

The code runs at C-like speeds

Parallelisation is an integral part of the language

Non-application based development is typically interactive nowadays

Python has shown us how a dynamically typed, interpreted language can be easily used to solve our problems

Rapid prototyping and low barrier to entry are important language characteristics

But what happens when you need to scale things up?

2-language problem



Typical solutions involve using 2 languages

Python, matlab, mathematica, etc. for investigatory work and prototype development

C, Fortran, specialised solutions (or even just recoding in Cython) for performance implementations

This is a huge development overhead

Julia

- is dynamically typed, but with optional types
- built-in types are equivalent to user-defined types
- is JIT compiled using LLVM
- utilises dynamic multiple dispatch
- has full metaprogramming capabilities
- can call C and Fortran libraries natively

Dynamical language with optional Typing

Python is dynamically typed

C is statically typed

Knowing the type of a variable increases the speed of the compiled program

Needing to annotate your code with types slows down the coding process

So in Julia it's optional

Use of types is largely *not* required to obtain the performance boosts!

It will only help slightly

But *multiple dispatch* largely removes the overheads of type checking which exist in Python

It's mainly useful for ensuring that the correct code gets called in the right place

Dynamical language with optional Typing

Additional comments on types:

Despite variables having different types, promotion to compatible types automatically occurs

eg. `>> 1 + 0.234` `Int + Float = ?`

Most importantly, no meaningful difference between built-in data types (eg. `Int64`, `Float64`) and user-defined types

LLVM is a compiler infrastructure project

Julia generates LLVM Intermediate Representation (IR) code

LLVM takes this and optimises it, then compiles it for the host system

Finally the machine representation executes

This approach, apart from being fast, allows for inspection of the intermediate code

Multiple dispatch



Overloaded functions with late binding...

Class based OOP = dynamic single dispatch

```
A.function_call()
```

Function overloading (C++) = static multiple dispatch

```
set_parameter(A, 3.0)
```

Julia methods are defined and called based on acceptable parameter data types = dynamic multiple dispatch

```
my_func(A,B,C,22.3)
```

This feature is where much of the speed comes from!

Macros are applied at code interpretation-time

This allows for automated code generation (which makes your code cleaner to read)

And also, optimisations for the actual number of parameters or values used are possible (eg. explicit calculation of mathematical equations rather than generic solutions)

C and Fortran can be called ‘natively’



Many numerical libraries are written in C or Fortran

Julia can call these library functions directly:

```
ccall(:clock, Int32, ())
```

The JIT instructions generated are identical to those generated by a C compiler, no overhead!

Must be from a shared library

You are responsible for Type compatibility!

<http://julialang.org>

<http://docs.julialang.org>

<https://groups.google.com/forum/#!forum/julia-users>

Command-line invocation

REPL

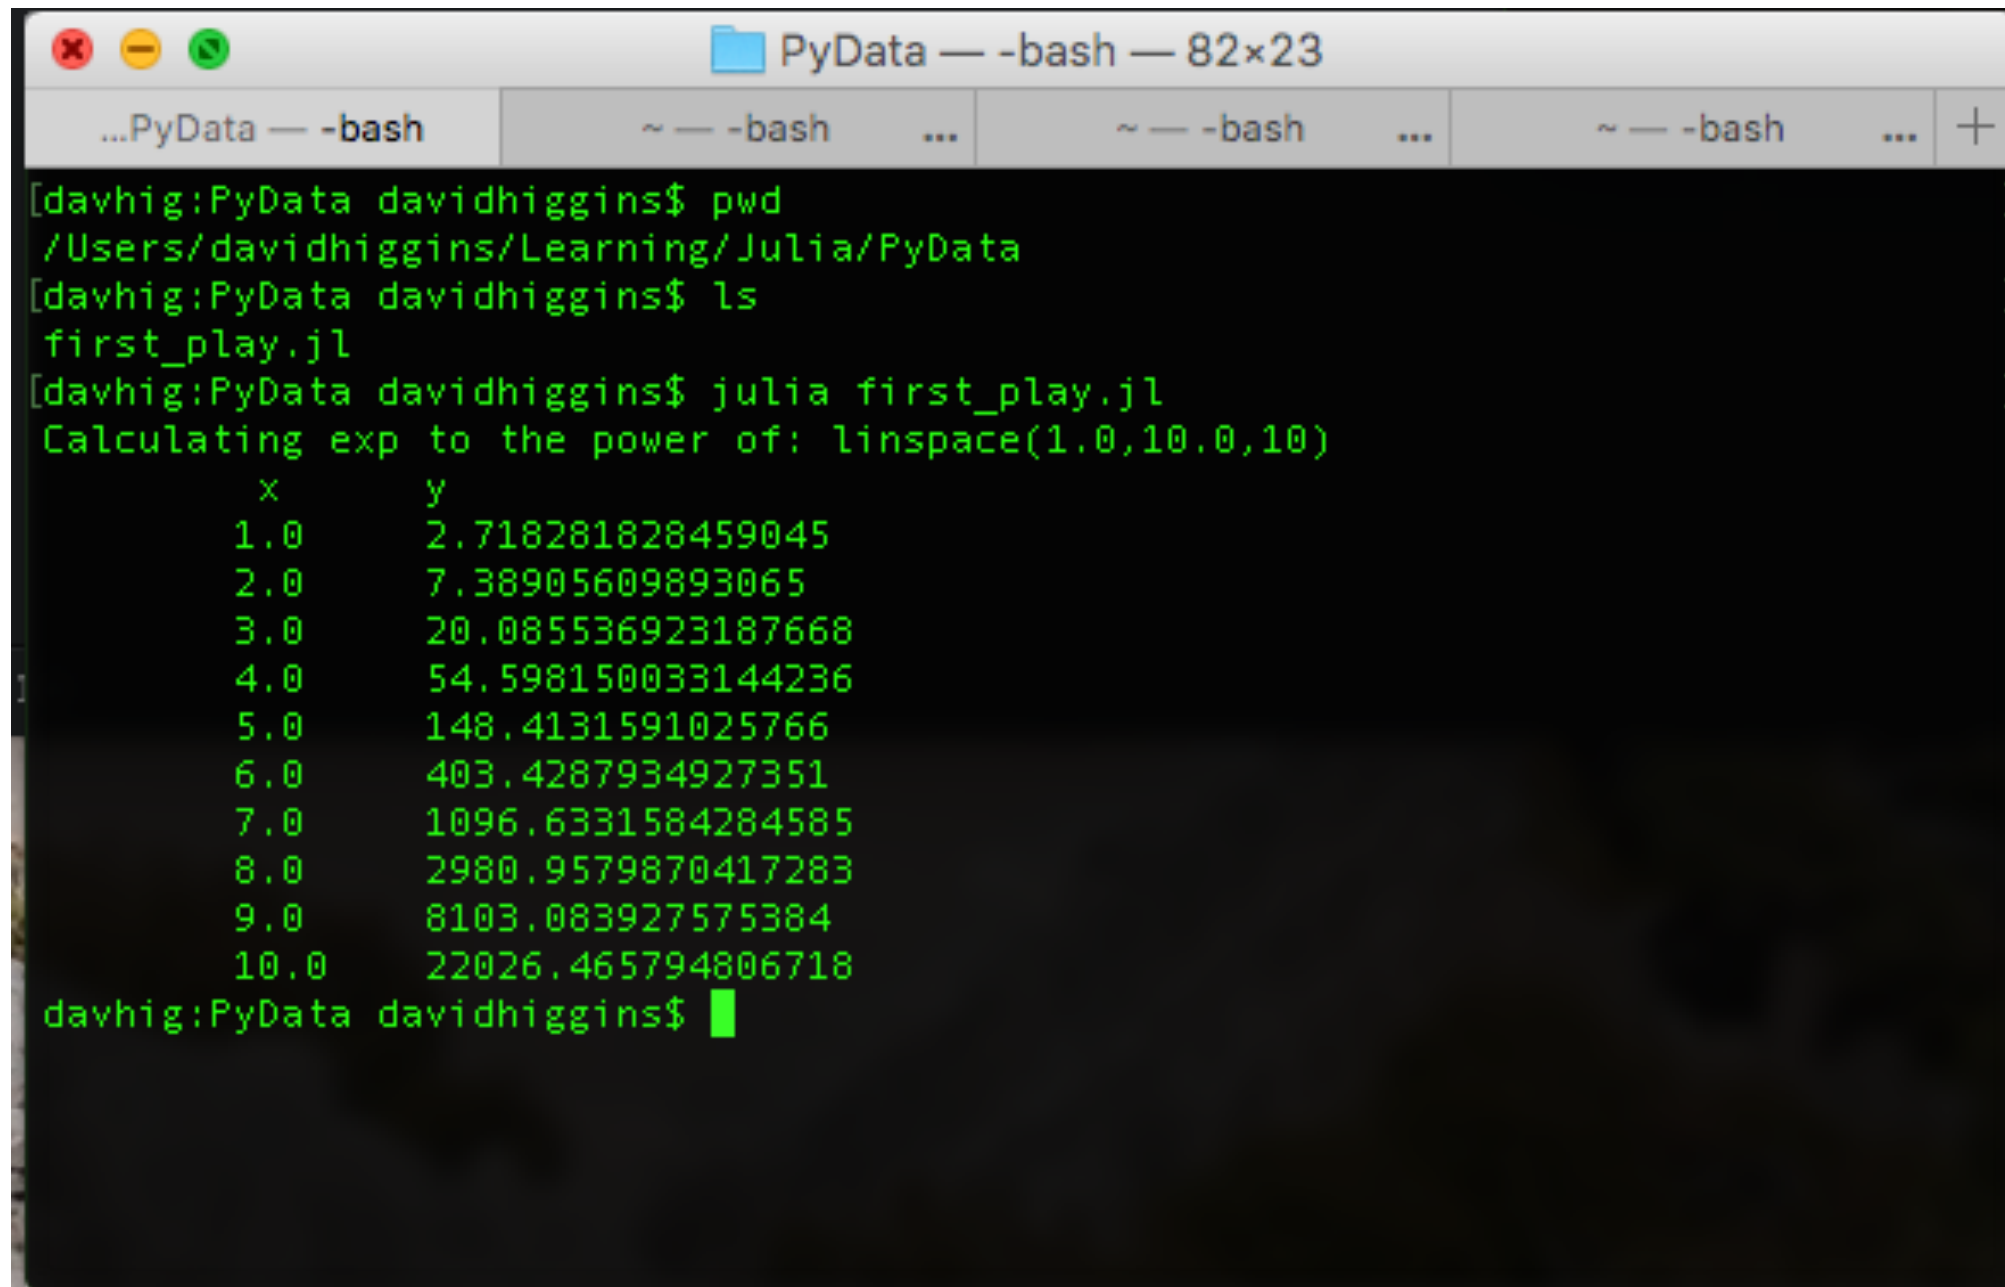
Jupyter (IPython like notebooks)

Juno + Atom (formerly with LightTable)

Interacting with Julia



Command-line invocation

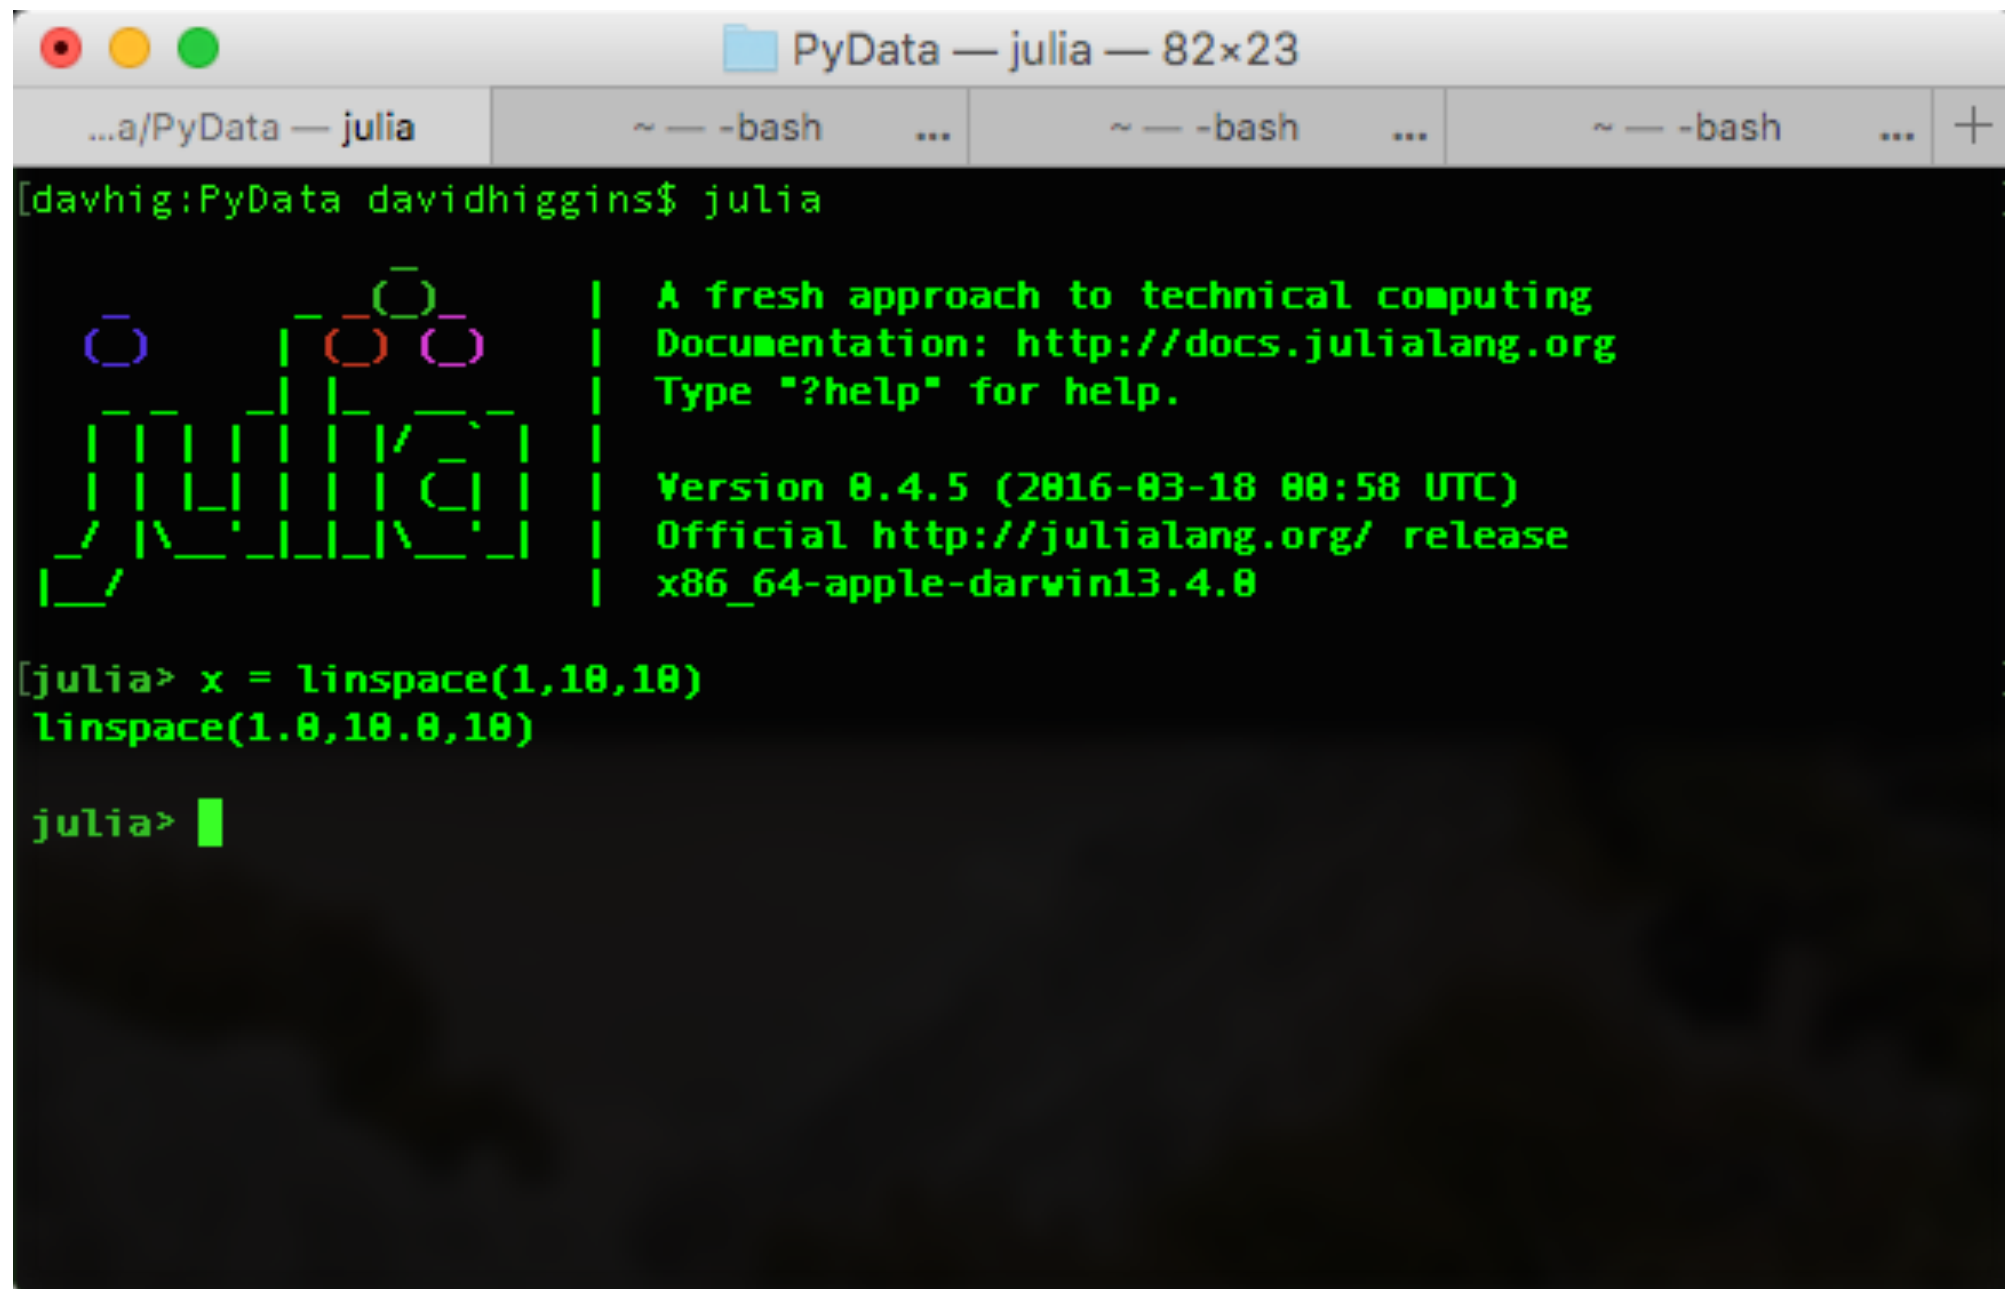
A screenshot of a terminal window titled "PyData — -bash — 82x23". The terminal shows a user named "davidhiggins" in a directory "PyData". The user runs "pwd" and "ls", then runs "julia first_play.jl". The output shows a calculation of the exponential function for values from 1.0 to 10.0.

```
[davhig:PyData davidhiggins$ pwd
/Users/davidhiggins/Learning/Julia/PyData
[davhig:PyData davidhiggins$ ls
first_play.jl
[davhig:PyData davidhiggins$ julia first_play.jl
Calculating exp to the power of: linspace(1.0,10.0,10)
      x      y
    1.0    2.718281828459045
    2.0    7.38905609893065
    3.0   20.085536923187668
    4.0   54.598150033144236
    5.0  148.4131591025766
    6.0  403.4287934927351
    7.0 1096.6331584284585
    8.0 2980.9579870417283
    9.0 8103.083927575384
   10.0 22026.465794806718
davhig:PyData davidhiggins$
```

Interacting with Julia



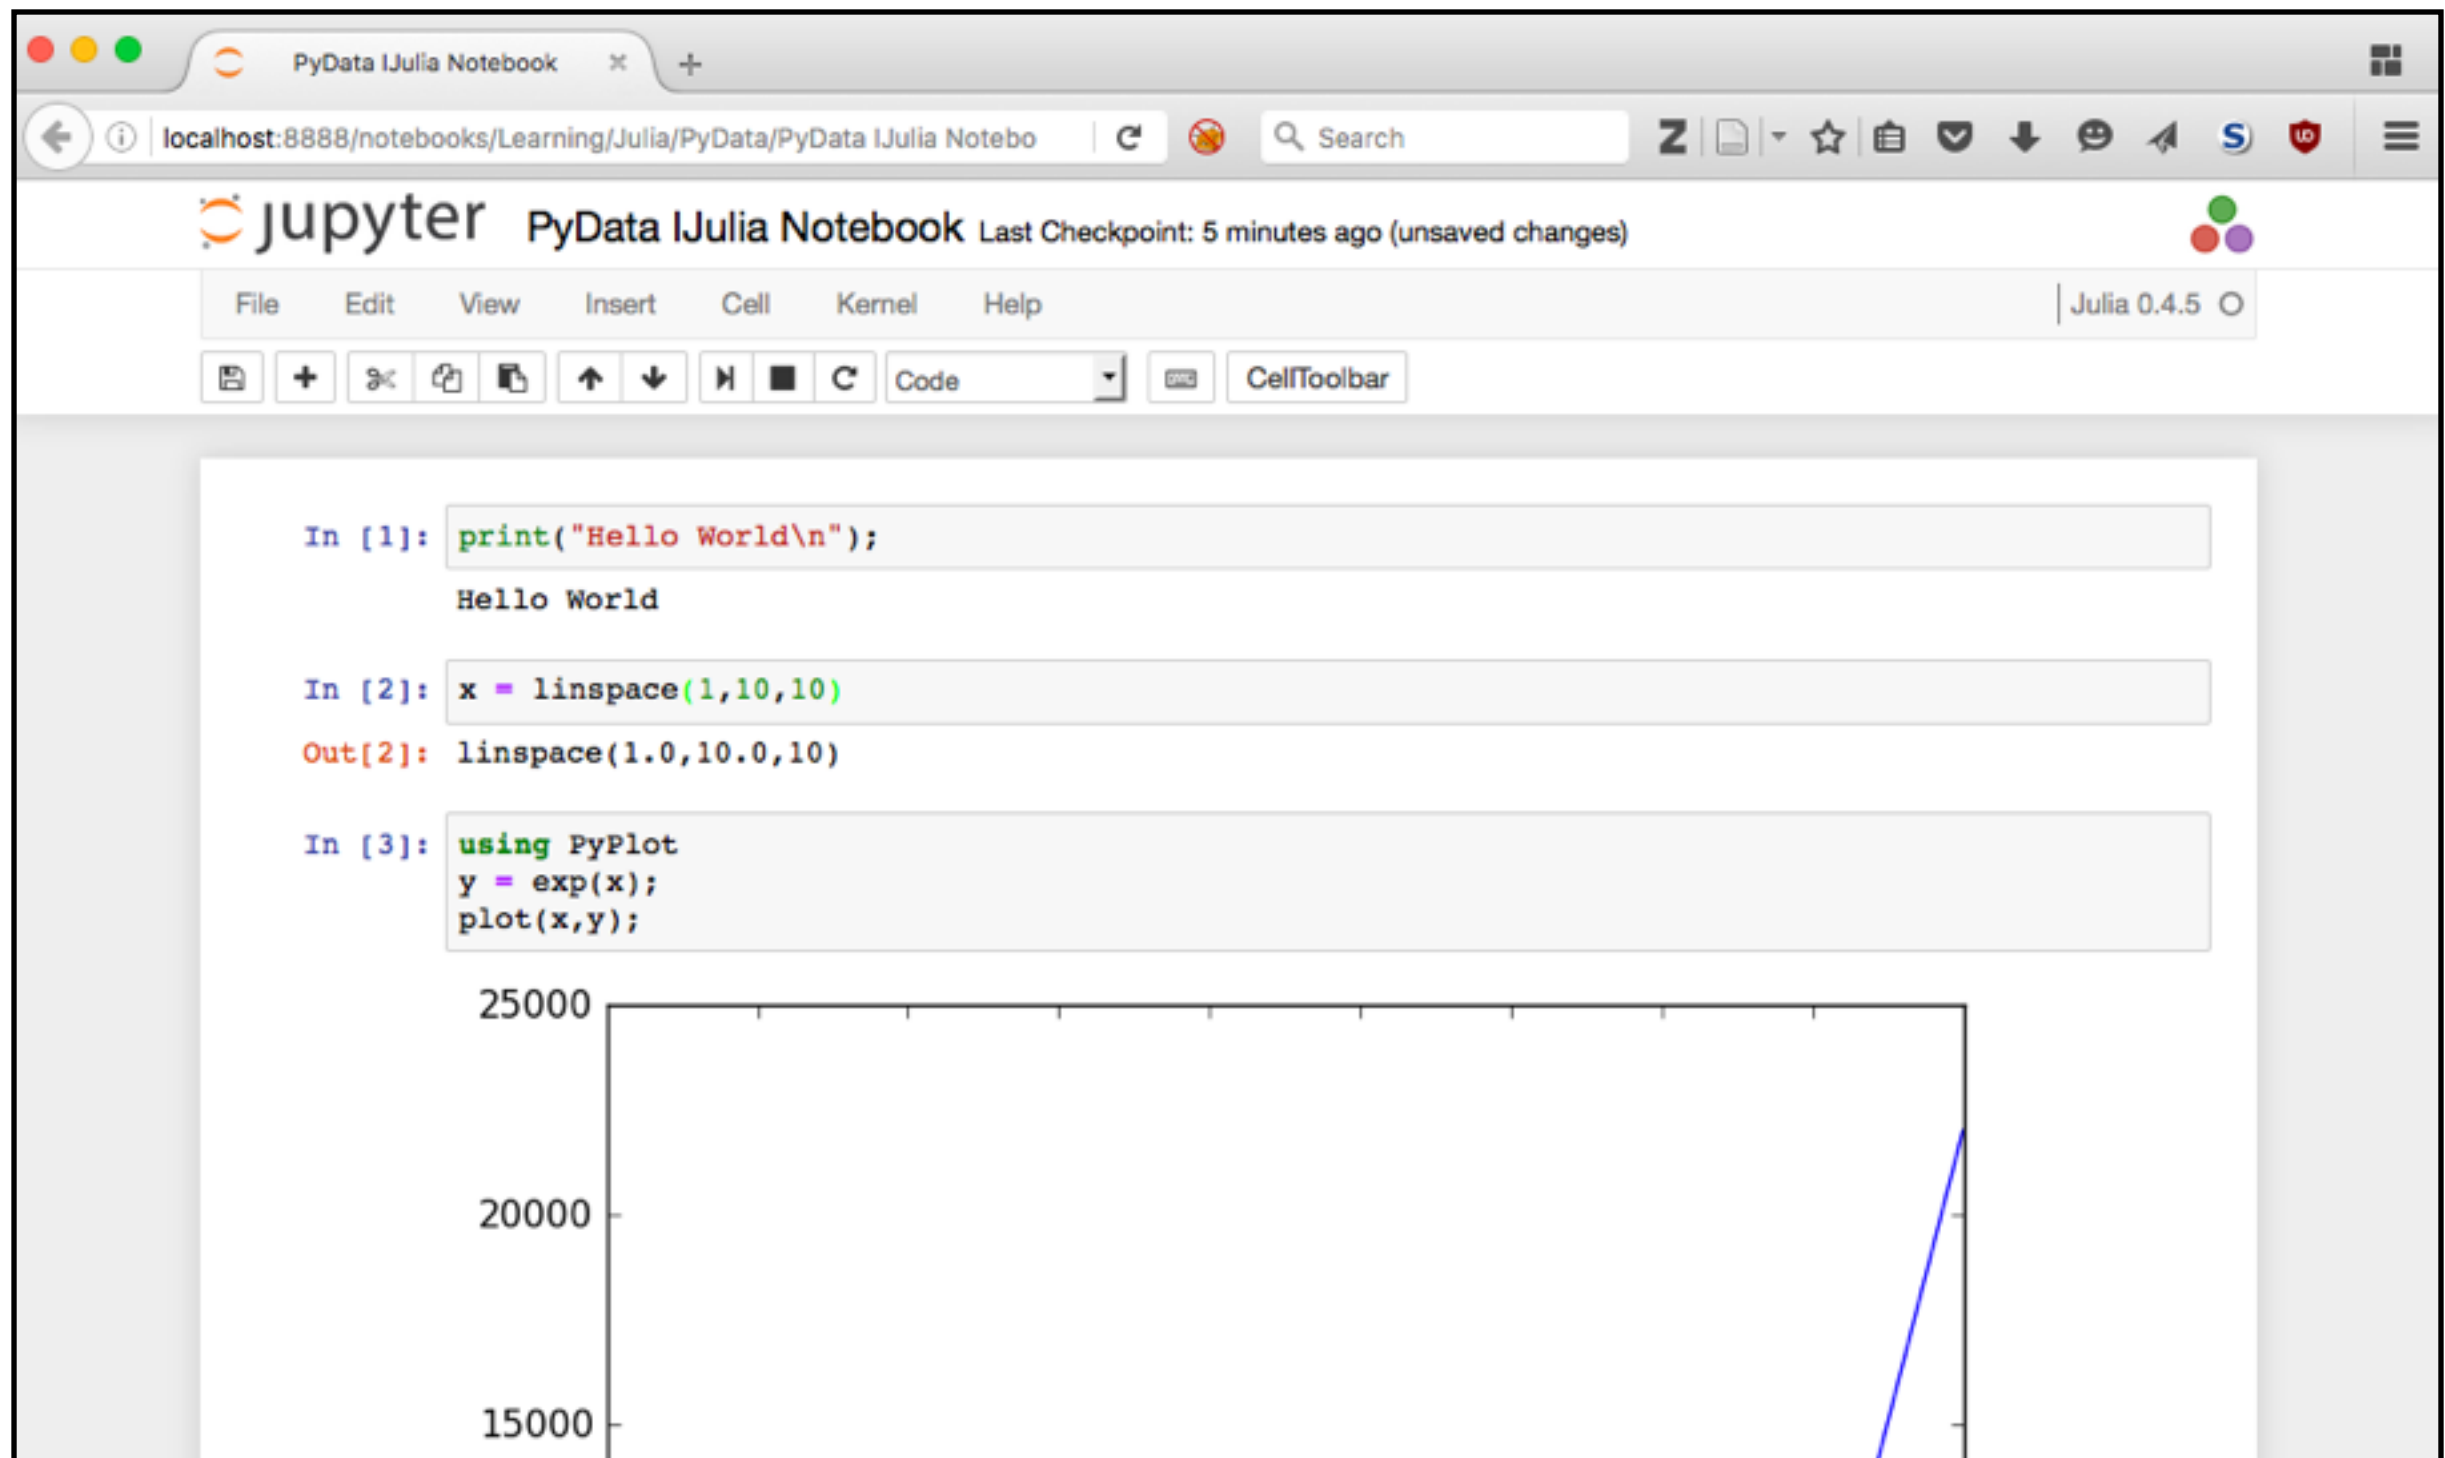
Read-Eval-Print-Loop (REPL)

A screenshot of a terminal window titled "PyData — julia — 82x23". The window shows the Julia REPL interface. The prompt is "[davhig:PyData davidhiggins\$ julia]". The output displays the Julia logo, which is a stylized grid of characters. To the right of the logo, the following text is shown: "A fresh approach to technical computing", "Documentation: <http://docs.julialang.org>", "Type \"?help\" for help.", "Version 0.4.5 (2016-03-18 00:58 UTC)", "Official <http://julialang.org/> release", and "x86_64-apple-darwin13.4.0". Below this, the user enters the command "[julia> x = linspace(1,10,10)]" and the output is "linspace(1.0,10.0,10)". The prompt "julia> " is shown at the bottom with a cursor.

Interacting with Julia



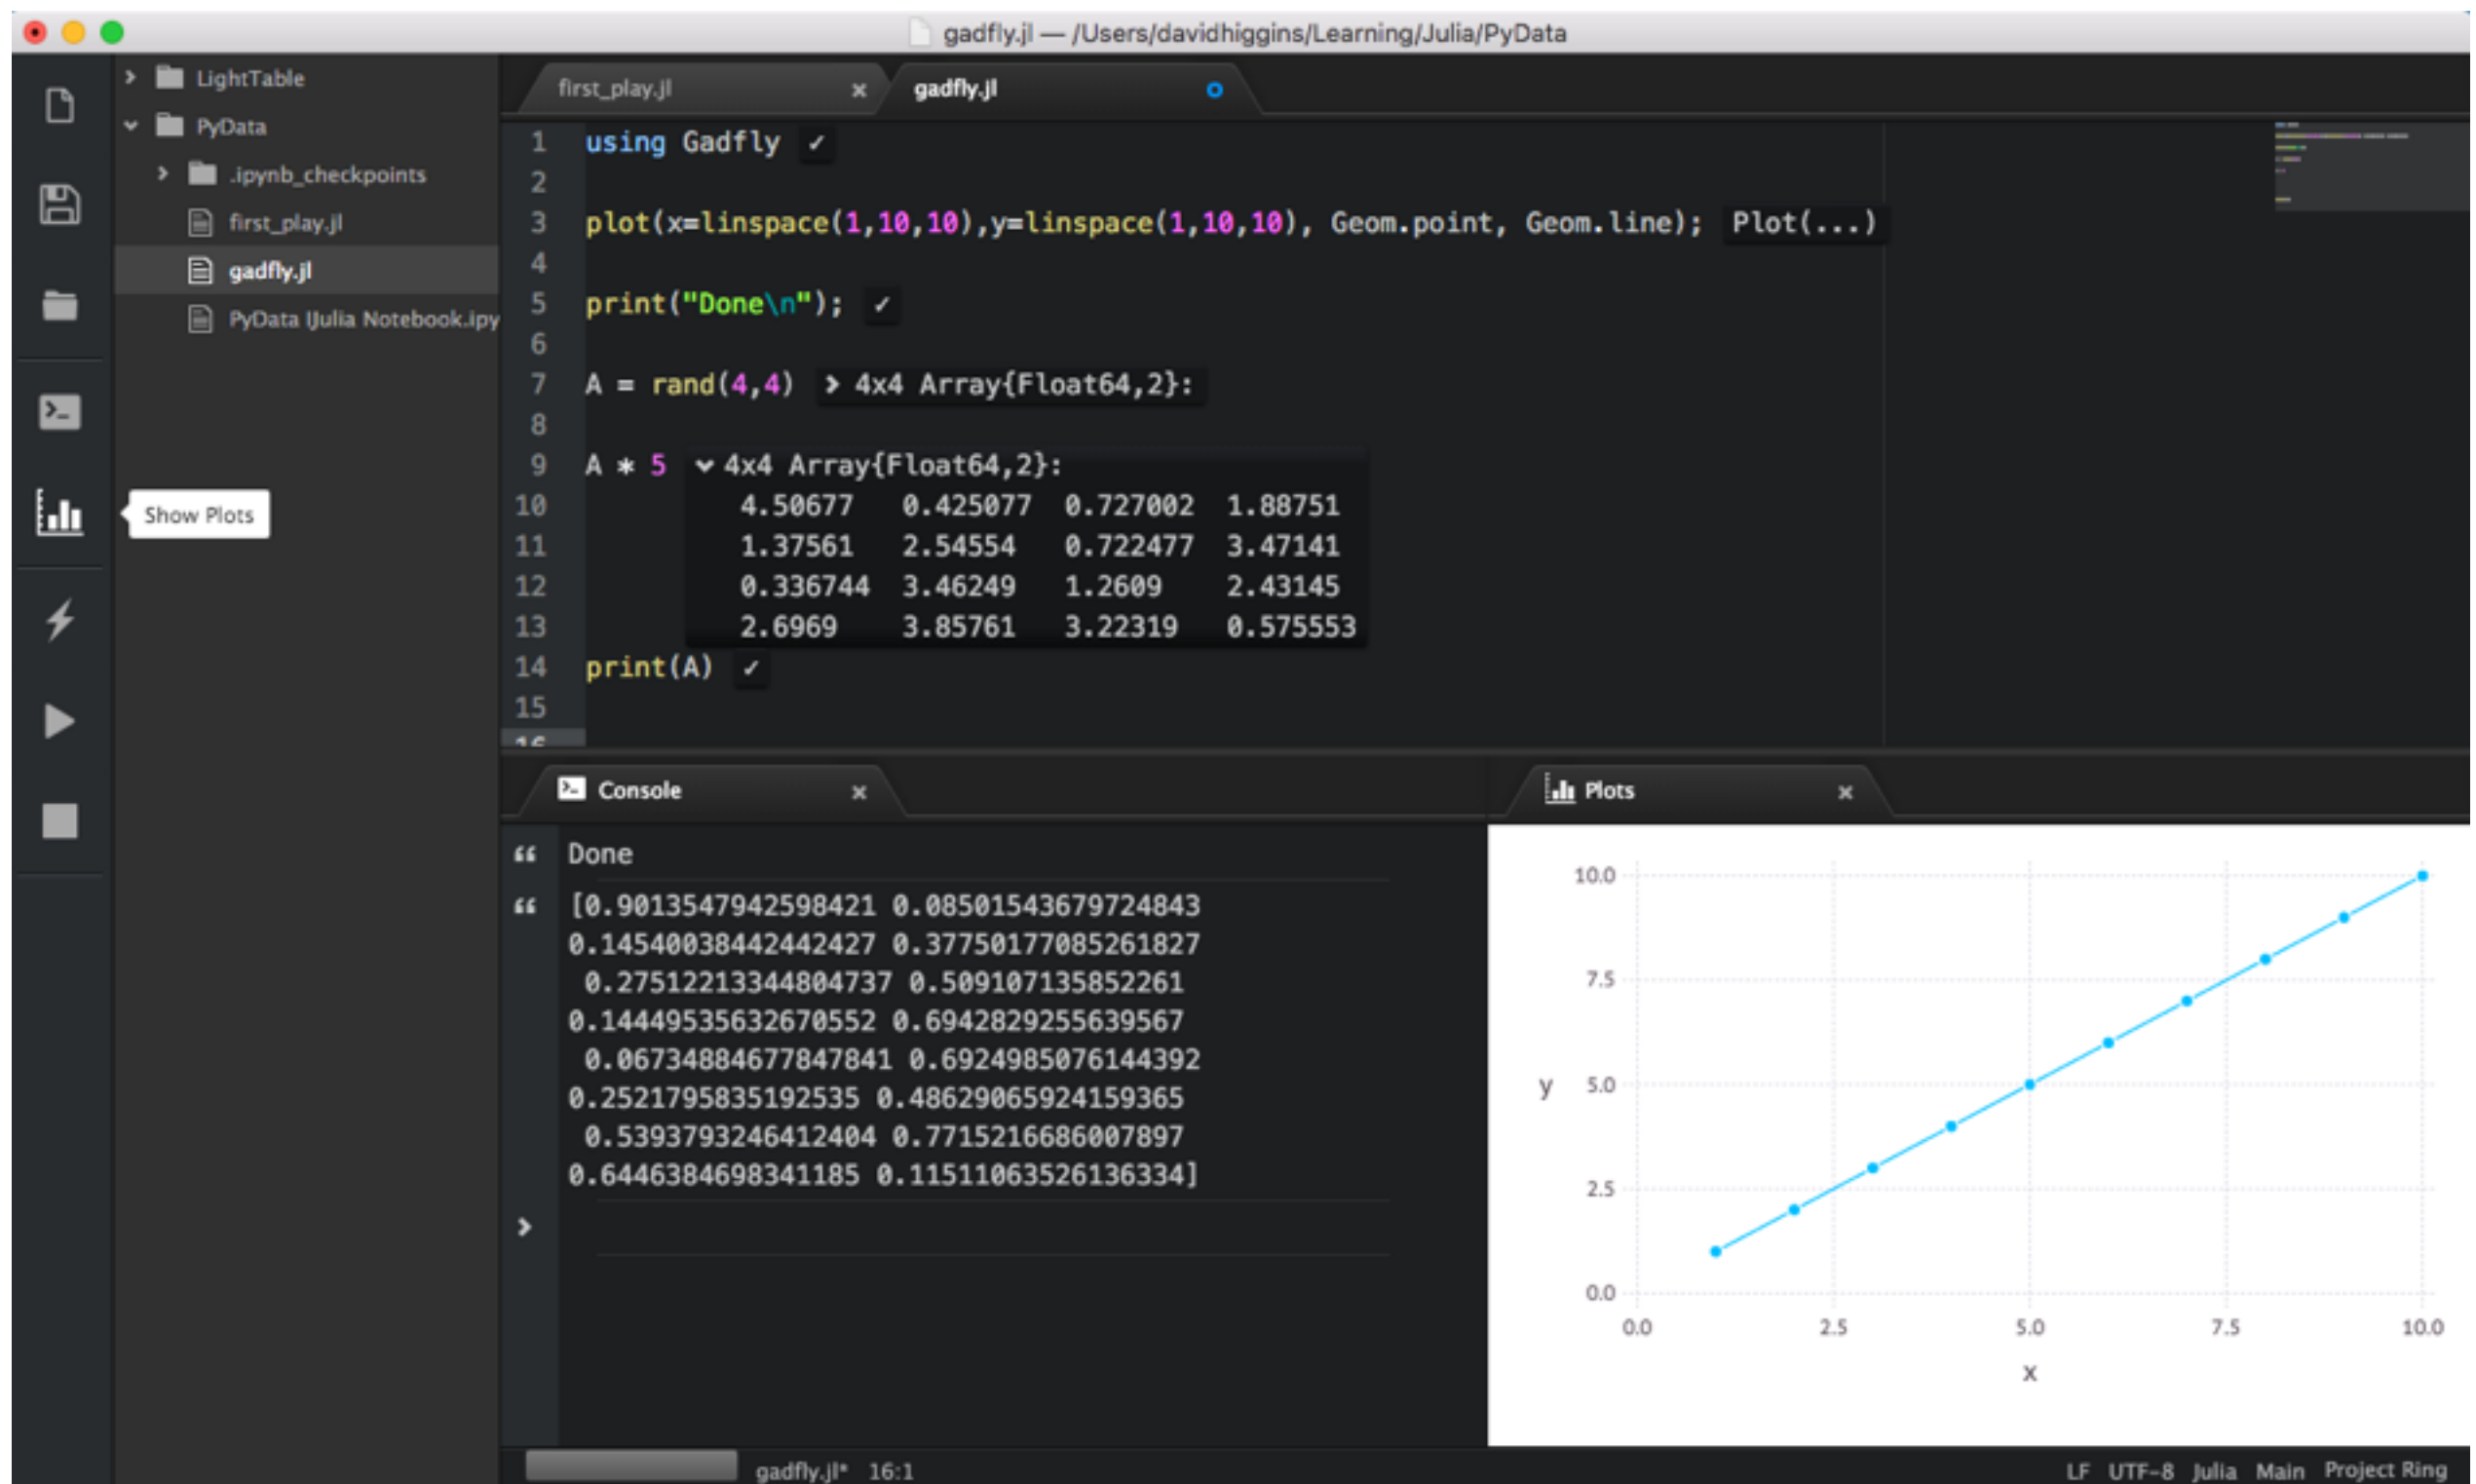
Jupyter Notebooks (IJulia)



Interacting with Julia



Atom + Juno (formerly used LightTable)



Forget about indentation



```
37 # these are the trial inputs
38 function generate_test_sequence(seq_length::Int64)
39     # linspace sequence:
40     #x = linspace(-1,1,seq_length);
41
42     # points are uniform randomly distributed:
43     if(use_cts_random_inputs)
44         x = rand(Uniform(problem_left_bound,problem_right_bound), seq_length);
45     end
46
47     # alternating +/-1 sequence
48     if(use_binary_alternating_inputs)
49         x = zeros(seq_length,1);
50         for(i=1:seq_length)
51             x[i] = -(-1)^i;
52         end
53     end
```


Switch to Jupyter notebook

Julia in the numbers



493 contributors (this week)

966 registered packages



[Bezanson; ArXiV]

Julia is

- quick to write
- fast to run
- a general purpose programming language
- designed from the ground up for numerical computing
- with parallel processing
- and easy scalability
- all of your existing tools still work in it

Thank you



Organisers

Audience

Language designers and developers

<http://julialang.org>

<http://docs.julialang.org>

<https://groups.google.com/forum/#!forum/julia-users>